


BLACKBELT ON AGENTIC ABM — MEGABRAIN ACADEMY — MODULE 4

4. ABM Metrics

Three groups of metrics, the waterfall funnel by tier, and how to prove progress before deals close.

4.1 The 3R framework: Revenue, Relations, Reputation



Before you can prove an ABM program works, agree what "working" means. Three groups of metrics, tracked together — not revenue alone.

WHY THIS MATTERS


Stop counting leads. Start counting companies.

"We sign contracts with COMPANIES" — a line from the source training, aimed at a room full of marketers still reporting MQLs and click-through rate. David Ogilvy, decades earlier, made the same point: don't count the people you reach — reach the people who count.

DAYS USING MEDIAN ▾



START RENDER VS BOUNCE RATE




SESSIONS



Case study: JLL's pilot ABM program, plan vs. actual


JLL pilot ABM program: plan vs. actual, first 2 quarters



13 target accounts, a demand map worth \$21.2M across 49 buying centers, 144 validated DM/influencer contacts. Six months in, actuals beat plan on every tracked metric. (Contact-map detail anonymized here — the original named real individuals at a real account.)


4.2 Reputation — the trust metrics

Customer loyalty ladder: from Prospect to Engaged



A share-of-wallet framework for judging whether a target account is actually deepening its relationship with you, not just buying once.

Champion vs. Influencer vs. Coach



A real champion has all three: incentive (something personal riding on your deal), influence (social capital to move others), and information (they help you navigate the account). Missing incentive = an Influencer. Missing influence = a Coach.

Four tests to confirm you actually have a champion


- Channel switch — end a call by asking to move to a more personal channel (text, WhatsApp); a real champion says yes without hesitation
- Ask them to describe their own buying process — a committed champion maps it out freely; a hesitant one keeps the map to themselves
- The small ask — request an intro or five minutes of feedback on a deck; did they actually do it?
- Teach-back — ask how they'd explain your solution internally with you not in the room, and offer to role-play the objections they'd get

Reputation — possible metrics

- Number of coaches / champions developed per target account
- Awareness of the target account's internal strategic initiatives
- Brand-perception trend and target-account website traffic
- Opt-in rate: downloads of thought-leadership content
- Campaign response rate; satisfaction / NPS
- Customer loyalty trend (ladder stage); references given; case-study consent


4.3 Relations — coverage & engagement

Weighting engagement: an "attribution in minutes" model




Assign every touch type a weight, then sum weighted minutes per contact and per account. Two published examples of the same idea: Grant Thornton's points-based engagement score, and Localytics' engagement-minutes $\times$ role-multiplier model (5x for CXO down to 1x for "other").

Illustrative engagement heat map: department × seniority



The same weighted-minutes data, rolled up by department and seniority level, shows you exactly where in the buying committee you're actually landing — and where you have zero coverage.

Account engagement over time, against an MQA threshold



MQA = Marketing Qualified Account: the account has earned enough weighted engagement to justify a hand-off to sales' most experienced closers, not a generic SDR sequence.

The "dark funnel": what attribution never sees

- Chris Walker's core argument: traditional b2b attribution models lose accuracy as more of the real buying journey moves outside trackable channels
- Untrackable signals that still move a deal: competitor research, anonymous site visits, word-of-mouth, internal brainstorming, community discussions, online reviews, a colleague's Slack message
- Practical response: don't over-trust last-touch or first-touch attribution alone — triangulate with self-reported "how did you hear about us" and directional signals (branded search volume, direct traffic) instead of chasing false precision


Relations — possible metrics

- New target-account contacts added to the database; contacts at DM/VP-and-above level
- Engaged contacts / engaged directors; meetings, demos, audits, feasibility studies held
- Directors participating in advisory councils or communities; joint innovation days
- Subjective engagement-strength score per contact; growth in touches with target-account DMs
- Attribution-weighted minutes and MQA status per account

4.4 Revenue — the funnel


Target accounts convert differently — same funnel, different economics

Target accounts vs. non-target accounts, same funnel



Illustrative comparison, adapted from the source material: current target accounts out-perform both brand-new target accounts and non-target accounts on every revenue metric that matters — bigger deals, higher win rate, shorter cycles, better retention, higher NPS.


The waterfall funnel: cohorts moving through the pipeline by tier



A simple goal-setting template: track how many accounts sit in each pipeline stage today, and where you expect that stack to shift over the next 1-3 months. This is the chart to rebuild in your own KPI workbook.

Illustrative dashboard: what a live ABM scorecard shows

Illustrative ABM dashboard mock-up



The source deck's real screenshots showed a live client's actual named target accounts inside their CRM — not reproduced here. This mock-up shows the same shape: coverage, cumulative engagement, and pipeline value in one view.

Revenue — possible metrics

- RFP / tender invitations from target accounts
- Number, growth rate and velocity of opportunities in the pipeline
- New pipeline opportunities this period; portfolio penetration (products sold per account)
- Stage-to-stage conversion rate and average sales-cycle length
- Win rate and average contract value
- Target-account cohort performance vs. a non-target-account control group

4.5 Proving progress before deals close

What "proof of progress" looks like, by ABM tier

	MONTHS 1-3	MONTHS 4-6	MONTH 6+ (success metric)
STRATEGIC 1:1	Charter signed, account research done	Workshop run, thought- leadership content built	Exec-level meetings with key DMs
ABM LITE (cluster)	Plan approved, "sold" internally	Content distributed via creative touchpoints	Deals logged in CRM
PROGRAMMATIC	Campaign + comms materials built	Nurture & engagement live	Closed deals, revenue booked


A signed contract can be 6-12+ months out, especially at the 1:1 strategic tier. Each tier has an earlier, honest proxy for "this is working" that you can report on before revenue shows up.

Keep it simple: one KPI table is enough

Metric	Now	+3 months	+6 months	+12 months
Contacts in DB with email	10	20	25	30
Engaged contacts	3	10	13	15
Director-level meetings	—	3	6	—
Deals logged in CRM	—	1	2	—
Contract value (illustrative)	—	—	\$300,000	\$600,000

4.6 Agentic ABM in practice


MegaBrain: an agent workflow for this module



One weekly prompt, routed through MegaBrain's model-agnostic gateway, turns raw CRM and engagement exports into a written pipeline-health brief — which accounts moved, which stalled, and what to do about it — without a dedicated BI hire.

Other agent tools for metrics & reporting

Dreamdata — B2B revenue-attribution platform, now shipping AI agents that turn multi-touch attribution data into plain-language pipeline narratives. Databox — connects 130+ data sources into real-time KPI dashboards with an AI layer that surfaces summaries and forecasts. HubSpot Breeze Assistant — CRM-native AI that drafts quick-look reports straight out of your existing HubSpot data. Make / n8n — a scheduled workflow that pulls the week's funnel movement and posts a digest to Slack every Monday morning.



Exercises

Homework 1 — build your own 3R metrics plan

- Pick the ABM tier you're actually running (Strategic 1:1, ABM Lite/cluster, or Programmatic) and state your target account count for this program
- For each of the three groups — Revenue, Relations, Reputation — pick 2-3 metrics from this module's "possible metrics" lists that you can actually measure with your current tools
- Set a target for each metric at four horizons: now, +3 months, +6 months, +12 months — round numbers are fine, precision isn't the point yet
- Fill in the KPI workbook that ships with this module using your own numbers, not the illustrative example
- Share the filled-in workbook with whoever owns sales for these accounts before the next session — a KPI plan nobody on the sales side has seen isn't a shared plan

Homework 2 — build an attribution-in-minutes model

- List every touch type your team actually uses today: calls, demos, webinars, direct mail, gated content downloads, email opens, meetings
- Assign each one a weight (minutes, points — the unit doesn't matter, the relative ranking does) using this module's example as a starting point
- Add a role multiplier: weight a CXO's engagement higher than an individual contributor's, the way Localytics' model does
- Run the model against 3-5 real target accounts using whatever CRM/analytics data you can pull manually this week
- Flag which accounts already look like an MQA (marketing-qualified account) by your own threshold, and send that shortlist to sales

Homework 3 — map your own "proof before close" timeline

- Using the tier-by-tier grid from section 4.5 as a template, write down what "progress" will concretely look like for your program at months 1-3, 4-6, and 6+
- Be specific: not "more engagement" but "3 workshop-style calls booked with VP-or-above contacts at 5 of our 8 target accounts"
- Share this timeline with your manager or executive sponsor now, before the program starts — it's much easier to get buy-in on what counts as progress before the numbers exist than after
- Calendar a check-in at each milestone date to compare actual vs. this plan, the same way the JLL case study did

Homework 4 — run your first MegaBrain reporting brief

- Export this week's account and opportunity data from your CRM (or use last week's homework's manual pull if you don't have full CRM access yet)
- Using your MegaBrain account from Module 1, send a prompt asking the model to summarize funnel movement, flag stalled accounts, and suggest next actions — see the 4.6 workflow example
- Compare the agent's brief against your own read of the same data: what did it get right, what did it miss or get wrong
- Keep a running note of which prompts produce a brief you'd actually forward to your manager unedited — you'll refine this further in later modules